

Project1

March 22, 2020

1 Project 1

2 Linear Regression

3 1) Explore Data

```
[56]: import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
from plotnine import *

import statsmodels.api as sm

from sklearn.linear_model import LogisticRegression # Logistic Regression Model
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler #Z-score variables
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import r2_score, mean_squared_error

from sklearn.model_selection import train_test_split # simple TT split cv
from sklearn.model_selection import KFold # k-fold cv
from sklearn.model_selection import LeaveOneOut #LOO cv
from sklearn.model_selection import cross_val_score # cross validation metrics
from sklearn.model_selection import cross_val_predict # cross validation metrics
```

```
[57]: data = "https://raw.githubusercontent.com/cmparlett/pelleriti/
↳CPSC392ParlettPelleriti/master/Data/diabetes2.csv"
```


```
[52]: #grabs data from online link
diabetes = pd.read_csv(data)
#reads data from csv file
diabetes.head()
#looks at the first 5 observations of the csv file
```

```
[52]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	


	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: (ggplot(diabetes, aes("Age"))
+ geom_histogram(aes(fill = "Age"))
+ labs(title = "Age Count", x = "Age", y = "Count"))
```



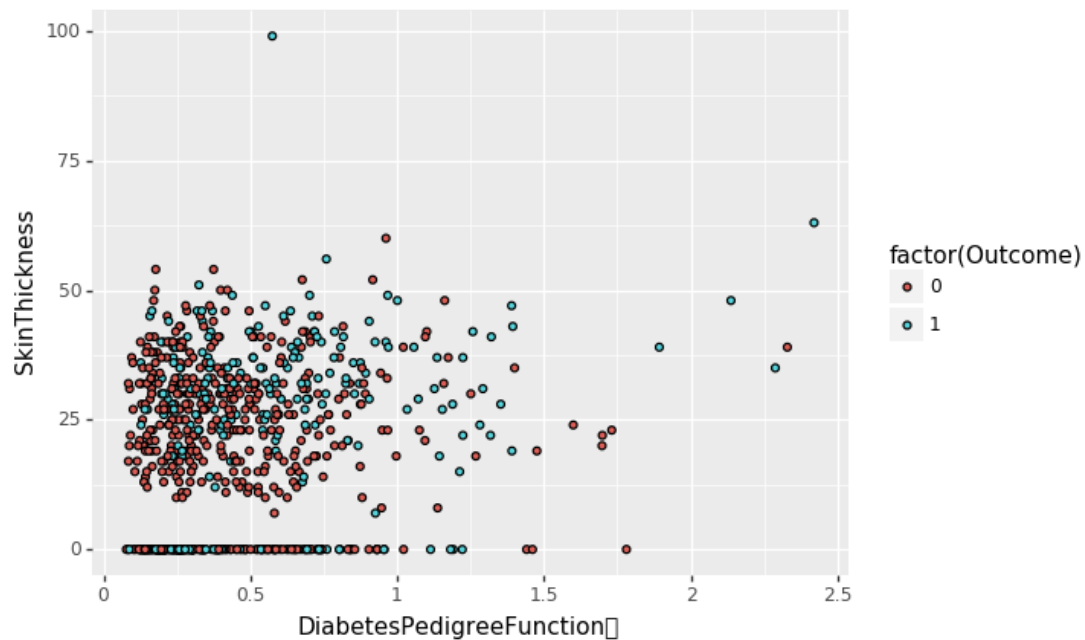
```
[4]: <ggplot: (301571917)>
```

```
[5]: (ggplot(diabetes, aes("Pregnancies", "SkinThickness"))
+ geom_point(aes(fill = "factor(Outcome)")))
```



```
[5]: <ggplot: (301578353)>
```

```
[6]: (ggplot(diabetes, aes("DiabetesPedigreeFunction", "SkinThickness"))
+ geom_point(aes(fill = "factor(Outcome)")))
```



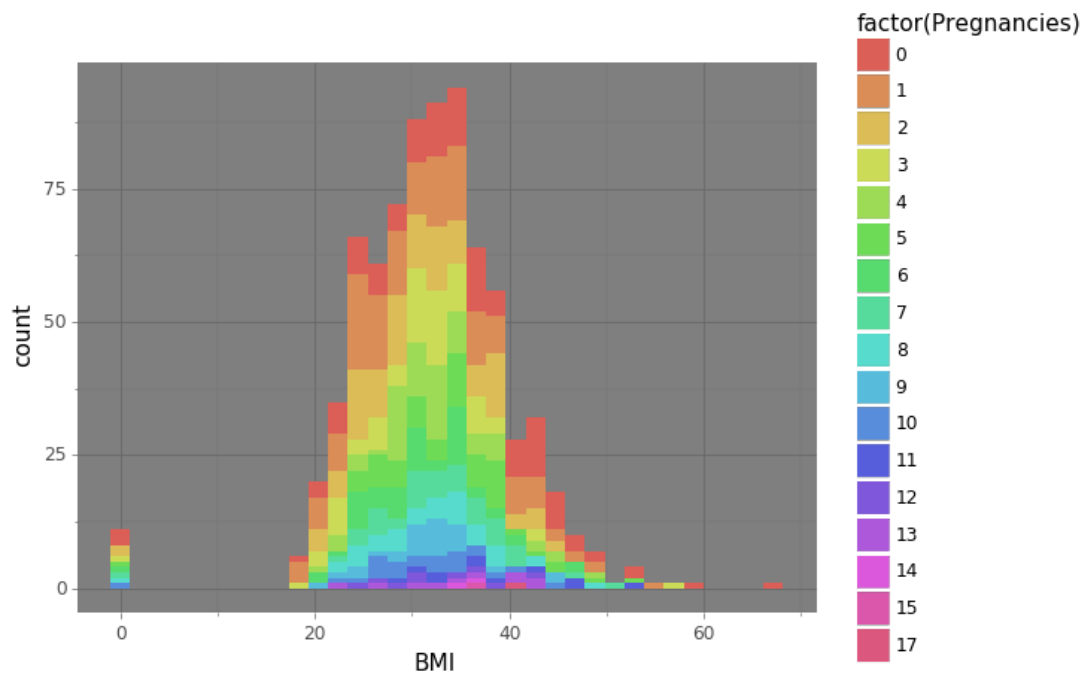
[6]: <ggplot: (301571853)>

```
[7]: (ggplot(diabetes, aes('Insulin', 'Glucose', color = 'factor(Pregnancies)'))  
+ geom_point() #creates point graph  
+ stat_smooth(method='lm') #smooths the slope line  
+ facet_wrap('~Pregnancies') #seperates each mfr into its own graph  
+ theme_minimal()  
+ labs(title = "Insulin and Glucose",  
x = "Glucose", y = "Insulin")) #labels the graph
```



[7]: <ggplot: (302875505)>

```
[8]: (ggplot(diabetes, aes('BMI'))  
+ theme_dark()  
+ geom_histogram(aes(fill = "factor(Pregnancies)")))
```



[8]: <ggplot: (302919769)>

4 2) Building My Model

[9]: `diabetes.head()`

```
[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[10]: # creates predictors
predictors = ["Glucose", "BloodPressure", "Insulin", "Age", "SkinThickness",
              ↪ "Pregnancies", "DiabetesPedigreeFunction"]
```

```
#creates test data and training data
X_train, X_test, y_train, y_test = train_test_split(diabetes[predictors],
↳diabetes["BMI"], test_size=0.2)
```

```
[11]: #Standardization
zscore = StandardScaler()
zscore.fit(X_train)
Xz_train = zscore.transform(X_train)
Xz_test = zscore.transform(X_test)
```

5 3) Evaluate your model

```
[12]: # create linearRegression model
LR_Model = LinearRegression()
```

```
[13]: # fit logModel
LR_Model.fit(Xz_train,y_train)
```

```
[13]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
[17]: BMI_preds = LR_Model.predict(Xz_test)
```

```
[18]: mean_squared_error(y_test,BMI_preds)
```

```
[18]: 47.737293995292326
```

```
[19]: r2_score(y_test,BMI_preds)
```

```
[19]: 0.23325950918080884
```

```
[19]: # ERROR
# My model did not do to well. My mean-squared error was 47.73

# ACCURACY
# Also my model had a r score of 0.23 which is not very good.
```

6 4) Interpret the coefficients to your model

```
[20]: coefficients = pd.DataFrame({"Coef":LR_Model.coef_,
                                "Name": predictors})
coefficients = coefficients.append({"Coef": LR_Model.intercept_,
                                "Name": "intercept"}, ignore_index = True)
```

```
[21]: coefficients
```

```
[21]:
```


	Coef	Name
0	1.546936	Glucose
1	1.120048	BloodPressure
2	-0.492824	Insulin
3	-0.227689	Age
4	3.038312	SkinThickness
5	0.017818	Pregnancies
6	0.377344	DiabetesPedigreeFunction
7	31.975407	intercept

```
[23]: # These coefficients show the affect of these variables on the BMI level
# For one stdv increase of Glucose there is a 1.547 increase in BMI
# For one stdv increase of BloodPressure there is a 1.120 increase in BMI
# For one stdv increase of Insulin there is a -0.493 decrease in BMI
# For one stdv increase of Age there is a -0.228 decrease in BMI
# For one stdv increase of SkinThickness there is a 3.038 increase in BMI
# For one stdv increase of Pregnancies there is a 0.018 in BMI
# For one stdv increase of DiabetesPedigreeFunction there is a 0.377 increaes
→ in BMI
# If there were no variables involved, y-intercept is 31.975
```

7 Logistic Regression

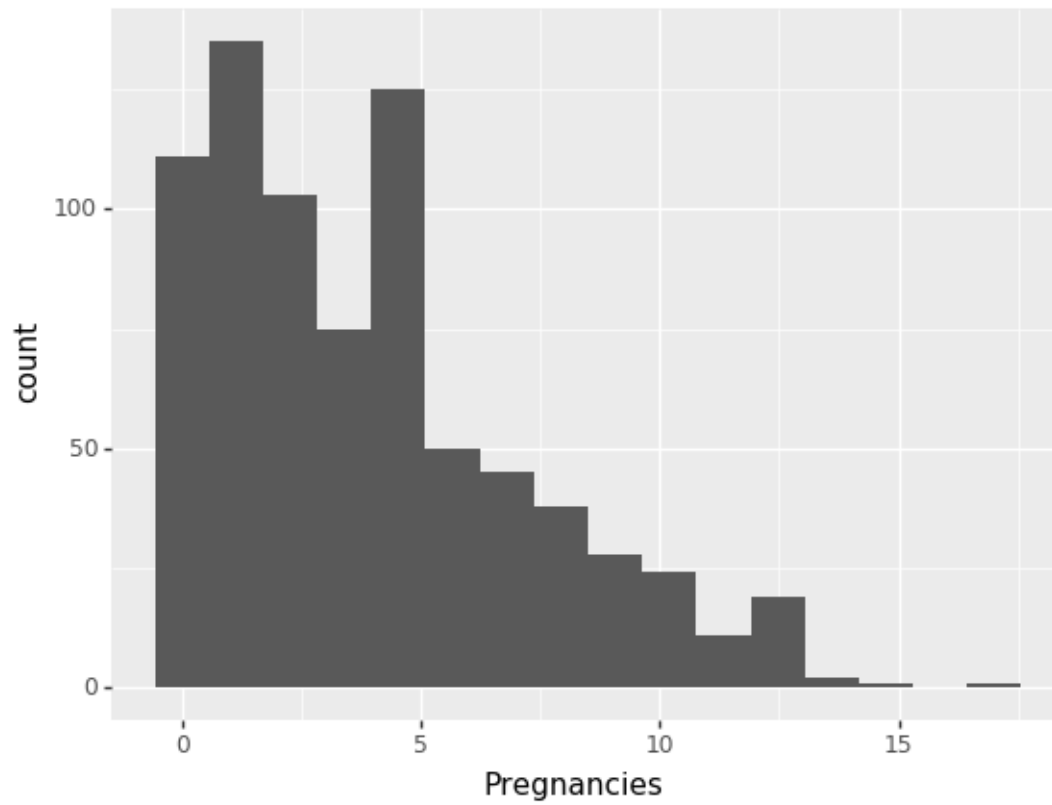
8 1) Explore Data

```
[49]: (ggplot(diabetes, aes("Glucose"))
+ geom_histogram())
```




[49]: <ggplot: (317381793)>

```
[46]: (ggplot(diabetes, aes("Pregnancies"))  
      + geom_histogram())
```

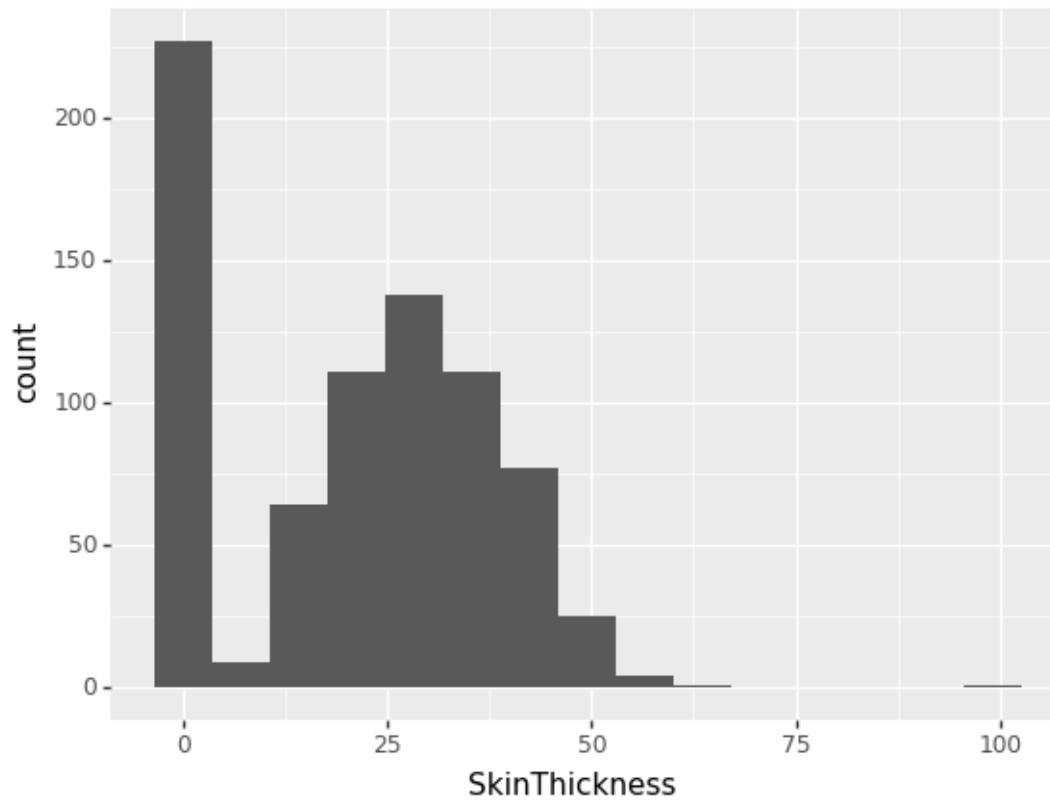
[46]: <ggplot: (317578725)>

```
[53]: (ggplot(diabetes, aes("Age", "BMI"))  
      + geom_point())
```



[53]: <ggplot: (317043029)>

```
[54]: (ggplot(diabetes, aes("SkinThickness"))  
      + geom_histogram())
```



```
[54]: <ggplot: (316943769)>
```

9 2) Building My Model

```
[53]: diabetes.head()
```

```
[53]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[63]: # Kfold

X = diabetes[["Glucose", "Pregnancies", "BloodPressure", "SkinThickness",
↳"Insulin", "BMI", "DiabetesPedigreeFunction", "Age"]]
y = diabetes["Outcome"]

# create k-fold object
kf = KFold(n_splits = 8)
kf.split(X)

# standardization
zScore = StandardScaler()
zScore.fit(X)
Xz = zScore.transform(X)

lr = LogisticRegression() #create model

acc = [] #create empty list to store accuracy for each fold
```

```
[64]: # Use a for loop to loop through each fold and train a model, then add the
↳accuracy to acc.

for train_indices, test_indices in kf.split(Xz):
    # Get your train/test for this fold
    X_train_k = X.iloc[train_indices]
    X_test_k = X.iloc[test_indices]
    y_train_k = y[train_indices]
    y_test_k = y[test_indices]

    # model
    model = lr.fit(X_train_k, y_train_k)
    # record accuracy
    acc.append(accuracy_score(y_test_k, model.predict(X_test_k)))

#print overall acc
print(acc)
np.mean(acc)
```

```
[0.7395833333333334, 0.8125, 0.7083333333333334, 0.75, 0.7708333333333334,
0.8125, 0.78125, 0.7916666666666666]
```

```
[64]: 0.7708333333333334
```

10 Evaluating my Model

```
[66]: len(diabetes)
```

```
[66]: 768
```


Metrics/Interpretation of model: I used the accuracy score to determine how accurate my model was. I got an accuracy score of 0.77 which is a pretty accurate model score so I am happy with my model

Justification for Cross-val technique: Since there are hundreds of observations in this dataset(768), I chose to do K fold as my cross-validation technique because it is less computationally expensive.

11 Data Viz

12 #1

```
[37]: (ggplot(diabetes, aes('BloodPressure', 'Glucose', color = "factor(Outcome)"))
+ geom_point() #creates point graph
+ stat_smooth(method='lm') #smooths the slope line
+ facet_wrap('~Outcome') #seperates each mfr into its own graph
+ theme_minimal()
+ labs(title = "BloodPressure and Glucose",
x = "BloodPressure", y = "Glucose")) #labels the graph
```



```
[37]: <ggplot: (317131621)>
```

13 #2

```
[59]: (ggplot(diabetes, aes('Age', 'DiabetesPedigreeFunction', color =_
  ↳ 'factor(Outcome)'))
+ geom_point() #creates point graph
+ theme_classic()
+ labs(title = "Age and DiabetesPedigreeFunction",
  x = "Age", y = "DiabetesPedigreeFunction")) #labels the graph
```



```
[59]: <ggplot: (317195853)>
```

```
[ ]:
```